Benchmark RADAR: Living Search Engine for Retrieval and Discovery of AI Benchmark Research

Koutian Wu[1,2] * Junjie Zhou[4] Ergan Shang[3] Jiayu Wang [5] Pengqian Han[6]

[1] Earth and Space AI [2] Tacite AI [3] Carnegie Mellon University

[4] Hangzhou Dianzi University [5] Xi'an Jiaotong University

[6] The University of Auckland

Abstract

Choosing an AI benchmark requires finding the test, its released artifacts, and evidence of how models perform on it. These records are scattered across papers, repositories, datasets, leaderboards, and model reports. Benchmark Radar brings them into a searchable catalog with source links, benchmark mentions, and evaluation settings attached to scores. At 2026-09-06, the catalog contained 1,283 benchmark records across four sources; 37 model reports supplied mentions of 110 benchmarks and 322 scores across 86 benchmark tracks. The reporting evidence reveals a concentrated set of widely used evaluations: GPQA Diamond appeared in 27 of the 37 reports. Eight benchmark tracks had a recorded score within five points of a 100-point ceiling, while six had later mentions at least 180 days after their latest archived score. These findings help readers identify evaluations with little remaining score headroom and gaps in score coverage. A worked prior-art search illustrates how a contributor used the catalog to assemble related work before designing a new evaluation. The system is available through a web dashboard, feeds, downloads, and local search clients.

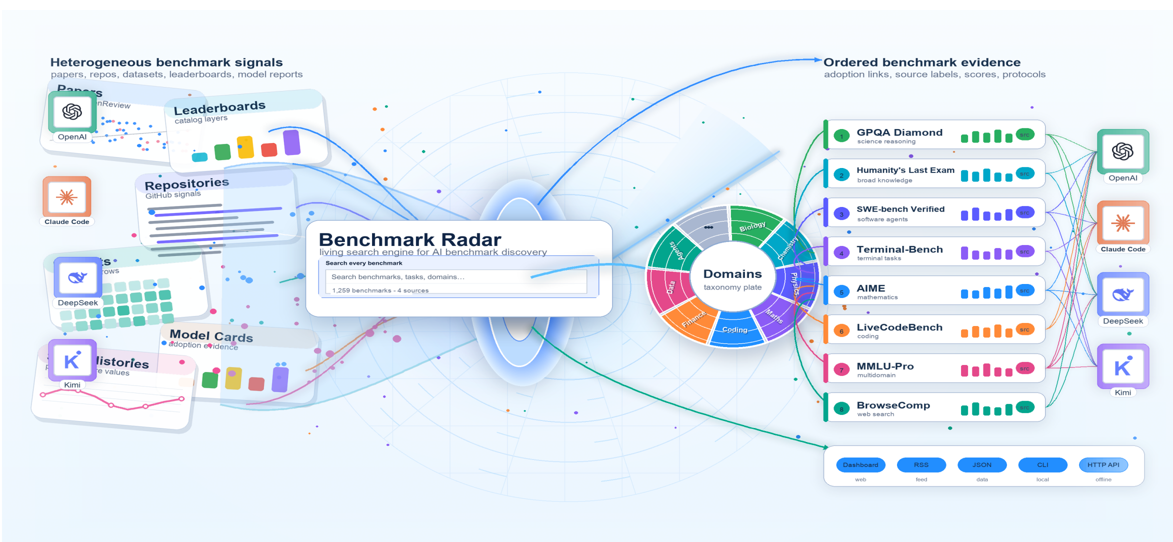

Figure 1: Benchmark Radar collects papers, repositories, datasets, and model reports, organizes them by topic, and makes them searchable alongside benchmark mentions and reported scores.

**Corresponding author: k@tacite.ai

1 Introduction

Before designing a new evaluation, a researcher needs to know which benchmarks already test the proposed capability, whether their data and code are available, and which models have been evaluated on them. Answering these questions requires following evidence across paper servers, code repositories, dataset hubs, leaderboards, and model reports. A benchmark name alone rarely supplies the test version, released artifacts, and evaluation settings needed to judge whether it fits the research question.

The range of evaluation tasks makes this search difficult. MMLU [11], GPQA [34], and Humanity's Last Exam [30, 43] test knowledge and reasoning. SWE-bench tests issue resolution [16]; LiveCodeBench uses date-based code-generation tests [15]; and Terminal-Bench and Long-Horizon-Terminal-Bench evaluate command-line agents [21, 24]. Researchers must also find domain-specific tests, including FinanceBench [14], MATH [12], SciBench [36], LAB-Bench [20], and ChemBench [25]. Choosing among them requires inspecting tasks, scoring rules, and released artifacts.

LLM Stats, OpenCompass, and Artificial Analysis provide benchmark catalogs, evaluation platforms, and model comparisons [3, 22, 28]. Benchmark Radar builds on their records by adding daily discovery and model-report evidence to a shared search interface. A reader can retrieve a benchmark, follow its paper and artifact links, and inspect which reports mention it and under which settings they report scores. Each source record keeps its identity and citations so the reader can trace a result back to the evidence.

We present the collection and search methods behind this catalog and examine what the reporting evidence reveals. At the cutoff, GPQA Diamond appeared in reports from 11 organizations, eight benchmark tracks had recorded scores within five points of their ceiling, and six had a gap of at least 180 days between their latest score and a later mention. These observations connect benchmark selection to adoption, score headroom, and missing measurements. Appendix B follows a contributor who used Benchmark Radar to gather prior art for a proposed evaluation.[1]

2 Methodology

2.1 System Overview

Benchmark Radar adapts the daily collection approach of BuilderPulse, which answers a fixed question set from more than ten sources each morning [39]. It applies this approach to benchmark discovery, catalog search, and score histories. Readers can look for recent releases, search for a task or benchmark name, inspect mentions in model reports, and compare scores collected under matching evaluation settings. Table 1 lists these uses and the interfaces that support them.

Benchmark Radar converts source records into common fields for names, descriptions, identifiers, dates, and supporting documents. A source observation is a collected record; an artifact is a paper, repository, dataset, release, or page identified from those observations. The benchmark catalog retains one record per benchmark within each source, including records with missing measurements. Model reports describe model development and evaluation [26]. We measure adoption by documented benchmark mentions and record numeric scores separately.

2.2 Collection and Source Health

Connectors are programs that retrieve records from public sources. Each collection run searches a 48-hour window, records counts and errors by source, removes future-dated rows, and requires healthy

[1] Source code: https://github.com/ktwu01/benchmark-radar.

Interface	Reader question	Role in Benchmark Radar
Today	What appeared or changed?	Ranks new records and links the source evidence.
Search box	Which records match this name, task, or domain?	Matches names, phrases, and query words in the benchmark index.
Leaderboard	Which benchmarks do labs report?	Counts benchmark mentions across curated model-card documents.
Scores	Which reported scores can be compared?	Groups scores by benchmark version, metric, and evaluation settings.
Trends and map	Which topics and sources recur?	Records source coverage so changes in collection can be considered when comparing topics.
CLI and HTTP	Can an analyst reproduce search locally?	Use one query service and a shared JSON response format.

Table 1: Reader questions and the interfaces used to answer them.

responses from a configured set of core sources before publication. The core sources at the cutoff were arXiv [4], Hugging Face Hub [13], and GitHub Search [9]. Optional routes include GitHub organizations, Hugging Face Papers, Kaggle datasets [17], Zenodo [42], OpenReview [29], Semantic Scholar [19], GitHub Releases [8], OpenAlex [31], Brave Search [5], 24 institution-operated research or engineering feeds, and a separate collector of public attention through Hacker News search on Algolia [1]. All three core sources were healthy at the cutoff. Semantic Scholar was rate-limited, Brave Search lacked an API key, and OpenReview returned no eligible rows. Appendix A lists the healthy connectors and monitored feeds.

The collection run at the cutoff fetched 867 rows, reduced them to 837 candidate records after removing duplicates, classified 322 as eligible, and recommended 110 for readers to review. Two collections ran within the cutoff window; after merging, the daily page contained 322 records, including 110 recommendations. Published records retain the source URL, retrieval time, parser version, source identifier, timestamps, authors or organizations when supplied, and a SHA-256 checksum for verifying the stored record. Raw responses and credentials are excluded.

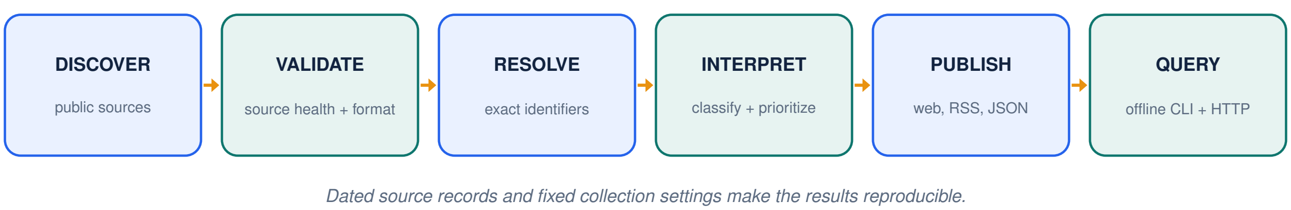

Figure 2: Collection and publication stages. Connectors retrieve public records and check source health and record formats. Matching DOI, arXiv, OpenReview, GitHub, or Hugging Face identifiers link observations to artifacts. Classification and priority scoring organize the records before publication. Web pages, feeds, and local clients expose the collected evidence.

Dated snapshots, record formats, parser versions, and checksums allow the published data to be rebuilt and checked. Publication requires healthy core sources; failures in optional sources remain visible. Installed clients validate an update before activating it. Identity resolution relies on exact identifiers and reviewed links between source records. The evaluated release passed 1,236 tests after a clean rebuild. Section 4 describes the remaining limitations.

2.3 Identity, Ranking, and Search

The resolver groups observations when they share a DOI, arXiv identifier, OpenReview identifier, GitHub repository, or Hugging Face artifact. Similar titles alone do not force a merge; they remain separate until review confirms that they describe the same artifact. This policy reduces false joins but leaves possible duplicates visible when papers, repositories, and catalogs lack a common identifier.

Search covers four catalog sources, as summarized in Table 2: LLM Stats, OpenCompass Hub, Artificial Analysis, and Model reports. The index contains 1,283 source records, including 110 model-report benchmarks. Of these model-report benchmarks, 86 have reported scores and 24 have no scores; both groups remain searchable. Search matches query words and phrases against record fields, including benchmark names. It searches all four sources together and preserves each result's source label. The same query against the same index produces the same ranking.

A separate priority score for daily recommendations combines relevance (35%), evidence (20%), recency (20%), and adoption (25%) on a 0–100 scale. These weights describe daily reading priority; search relevance depends on the query. The scoring configuration is included with the published data.

Catalog source	Rows	Primary use
LLM Stats	687	Benchmark names and descriptions with source provenance.
OpenCompass Hub	461	Paper, repository, release, and dataset links with source-specific identity fields.
Artificial Analysis	25	Current commercial evaluation catalog entries.
Model reports	110	Benchmarks mentioned in model reports, including records without scores.
Total	1,283	All four sources, retaining each record's source label.

Table 2: Catalog sources in the Benchmark Radar query index.

2.4 Adoption and Score Curation

Benchmark Radar records benchmark mentions and numeric scores from model-card and system-description documents [26]. This collection includes 37 model documents from 12 organizations. They mention 110 benchmarks whose names have been reconciled within the model-report collection. When a document contains an interpretable numeric value and sufficient context, the score archive records the benchmark version and metric (the instrument), along with evaluation settings (the protocol). These settings include examples supplied in the prompt (shots), tool access, attempts, and the evaluator when available. The archive contains 322 scores across 86 benchmark tracks. A score track groups results for a named benchmark; comparisons within it still require matching instruments and protocols. Scores from benchmark registries remain outside these model-report score histories unless they can be connected to a comparable protocol [7, 26].

3 Results

3.1 Coverage at the Cutoff

At the 2026-09-06 cutoff, the cumulative discovery collection held 10,702 source observations and 6,304 artifacts grouped by exact identifiers. Primary or structured sources supplied 10,473 observations, or 97.86% of the collection. This category covers scholarly records, institution-operated feeds, code repositories, and dataset or artifact hosts. Discovery observations, benchmark records, model-report mentions, and scores count different objects and must not be added together. Table 3 defines each unit.

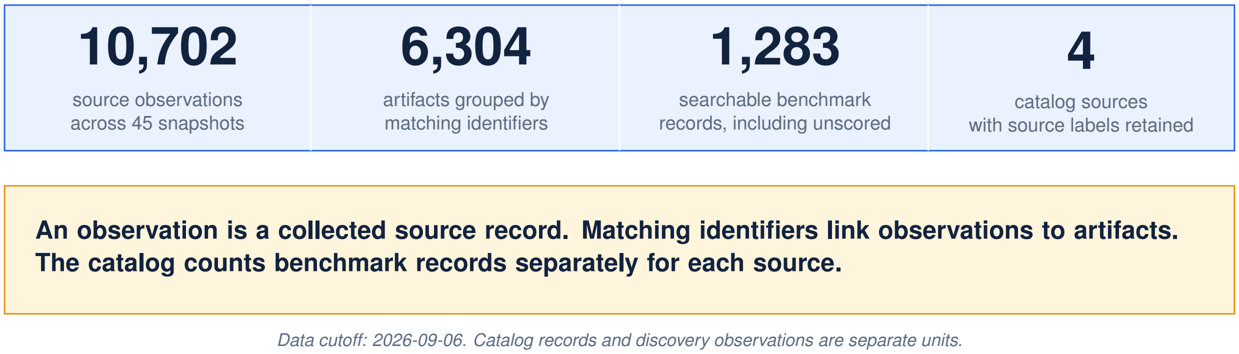

Figure 3: Source observations, resolved artifacts, searchable benchmark records, and catalog sources as of 2026-09-06.

Unit	Count	Interpretation
Snapshot	45	Dated collection records, including 41 observed and 4 simulated snapshots.
Source observation	10,702	Collected records with their source and retrieval information.
Unique artifact	6,304	Papers, repositories, datasets, releases, or pages grouped by matching identifiers.
Searchable entry	1,283	One benchmark record per source, including records without scores.
Adoption benchmark	110	Benchmarks mentioned in model reports, with names reconciled within that collection.
Model-card document	37	Curated reports from 12 organizations.
Curated score	322	Numeric results across 86 benchmark tracks.

Table 3: What each count measures at 2026-09-06.

3.2 Source Composition and Reliability

Five source labels account for 9,532 of 10,702 observations, or 89.1% of the cumulative collection. Because these sources supply most observations, their retrieval caps, rate limits, and availability can change the apparent topic mix. At the cutoff, the required sources were healthy, while optional routes varied in yield because of rate limits, missing credentials, or empty result windows.

Fifty-nine of 6,304 artifacts were linked across more than one source, or 0.94% of the artifact set. Matching exact identifiers reduces false joins; links between papers, code, and datasets still require review when identifiers are missing or inconsistent.

3.3 Relationship to Existing Catalogs

Benchmark Radar imports records from three benchmark registries: LLM Stats, OpenCompass Hub, and Artificial Analysis. These projects have different goals. LLM Stats publishes live benchmark rankings and benchmark-family pages [22]; OpenCompass is primarily an evaluation platform with a searchable dataset-statistics page [28]; and Artificial Analysis publishes model evaluations and a documented intelligence-benchmarking methodology [3].

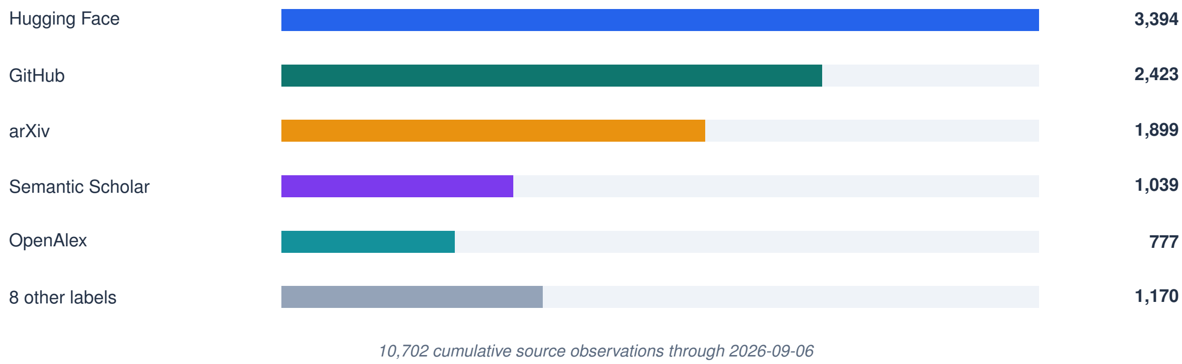

Figure 4: Discovery observations by source through 2026-09-06. The bars count collected observations, including repeat sightings, rather than distinct benchmarks.

The index contains 687 records from LLM Stats, 461 from OpenCompass Hub, 25 from Artificial Analysis, and 110 from Model reports, yielding 1,283 searchable source records. These counts describe each source's contribution to the catalog. Benchmark Radar adds daily discovery, links between artifacts with matching identifiers, and model-report mentions and scores. Readers can search this combined evidence while inspecting the source and definition of each record.

3.4 Benchmark Adoption in Model Reports

In the model-report collection, eight benchmarks appeared in documents from at least six organizations: GPQA Diamond [34], Humanity's Last Exam [30], SWE-bench Verified [16], Terminal-Bench [24], AIME [23], LiveCodeBench [15], MMLU-Pro [37], and BrowseComp [38]. GPQA Diamond appeared in 27 of 37 documents from 11 organizations, making it the most commonly mentioned benchmark in these reports. This concentration suggests that frontier-model reporting has partly converged on a common set of evaluations, even as researchers introduce benchmarks in other areas.

Score coverage lagged behind mentions for six benchmarks: GSM8K, HumanEval, MBPP, MMLU, DROP, and MATH-500. Each appeared in a model report at least 180 days after its latest archived numeric score; the observed gaps ranged from 343 to 775 days. Table 4 lists the dates. A reader investigating these benchmarks has newer reports to consult for score extraction.

Benchmark	Latest score date	Latest mention date	Gap (days)
GSM8K	2024-03-08	2026-04-22	775
HumanEval	2024-09-18	2026-04-22	581
MMLU	2025-01-22	2026-04-22	455
DROP	2025-01-22	2026-04-22	455
MBPP	2024-09-18	2025-11-20	428
MATH-500	2025-05-14	2026-04-22	343

Table 4: Benchmarks whose latest model-report mention follows their latest archived score by at least 180 days. Dates come from the score and model-report archives at the manuscript cutoff. The gap measures archive coverage.

3.5 Near-Ceiling Reported Scores

Eight of the 86 model-report benchmark tracks had a highest archived score within five points of a 100-point ceiling.[2] Table 5 retains the model and evaluation settings behind each value. Headroom is the metric's upper bound minus the reported score. For example, GLM-5.2 scored 99.2 on AIME 2026, leaving 0.8 points, while Gemini 3.1 Pro scored 99.3 on the Telecom domain of tau2-bench, leaving 0.7 points. These results identify test settings where further gains have little room on the reported scale.

Benchmark	Model and source	Score	Gap	Recorded settings
AIME 2026	GLM-5.2 [40]	99.2	0.8	Temperature 1.0; top-p 0.95; 163,840-token limit; GPT-5.5 (medium) judge.
Arena-Hard	Qwen3-235B-A22B [32]	95.6	4.4	Thinking mode; win rate.
DeepSearchQA	Kimi K3 [27]	95.0	5.0	F1; maximum reasoning effort.
HMMT Nov. 2025	GLM-5 [41]	96.9	3.1	Temperature 1.0; top-p 0.95; GPT-5.2 (medium) judge.
MATH-500	Qwen3-235B-A22B [32]	98.0	2.0	Thinking mode.
MathVision	Kimi K3 [27]	97.8	2.2	Maximum reasoning effort; Python available.
SWE-bench Verified	Claude Opus 5 [2]	96.0	4.0	Adaptive thinking; maximum effort; average of five trials.
tau2-bench, Telecom	Gemini 3.1 Pro [10]	99.3	0.7	High thinking level; Telecom domain.

Table 5: Highest archived scores for the eight near-ceiling model-report tracks. Scores use 0–100 scales; gaps are percentage points below 100. The metrics include accuracy, pass@1, F1, resolved-task rate, and win rate. Each row links to its reporting source.

The test version and settings determine what each small gap means. AIME 2026 and AIME 2024 use different problem sets; the 99.2 result above and the archive's older AIME 2024 result of 79.8 describe different tests. Likewise, the tau2-bench value above concerns Telecom, and the MathVision value includes Python access. Benchmark Radar preserves these details so readers can select the relevant test and inspect the score in context. Comparisons require checking the source documents for matching versions, reasoning budgets, tools, attempts, and evaluators [7, 26].

4 Limitations and Future Work

Lexical search can miss paraphrased benchmark names, renamed tasks, or descriptions of related ideas. Semantic retrieval could broaden the candidate set while preserving source links and retrieval explanations [33]. The worked example combines catalog queries with web search; evaluating retrieval quality requires relevance judgments over a broader set of queries and candidate records.

The collected evidence also has gaps. Public endpoints, credentials, and rate limits affect source coverage. Missing identifiers leave links between papers, code, and datasets unresolved. In the snapshot, 12,594 scores from benchmark registries remain outside the model-report score histories because they lack protocol fields such as prompt examples, tool access, attempts, evaluator, date-based test splits, or evaluation software. Matching archive labels alone cannot verify identical evaluation conditions or independent repeated runs. Further curation should capture these settings from primary reports and distinguish score-publication dates from evaluation dates.

[2] Junjie Zhou prepared the score-archive audit and report revisions underlying this analysis, documented in issue #457.

Task-level coverage is incomplete. The capability-classification component, KW-Bench, has 5,661 tracks without complete task-capability labels at the cutoff. A classification track is an artifact or a declared task within a benchmark suite, drawn from discovery observations. Completing these labels would help readers search by the ability a task tests.

5 Conclusion

Benchmark Radar makes 1,283 benchmark records searchable alongside their sources, model-report mentions, and score histories. The reporting evidence identifies widely used tests, near-ceiling results, and benchmarks whose score coverage has fallen behind later mentions. Researchers can use these records to assemble prior art, inspect evaluation settings, and target missing evidence for review. The next priorities are broader retrieval evaluation, task-capability labels, and more complete protocol capture from primary reports.

Author note. Koutian Wu led the work and manuscript preparation, built the initial collection pipeline, and implemented data aggregation across sources. Ergan Shang prepared figures and contributed copyediting. Pengqian Han contributed copyediting.

References

[1] Algolia. Hacker News Search API, 2026. URL https://hn.algolia.com/api. Accessed 2026-09-06.

[2] Anthropic. Claude Opus 5 system card, 2026. URL https://www.anthropic.com/claude-opus-5-system-card.

[3] Artificial Analysis. Artificial analysis intelligence benchmarking methodology, 2026. URL https://artificialanalysis.ai/methodology/intelligence-benchmarking. Accessed 2026-09-06.

[4] arXiv. arXiv API User Manual, 2026. URL https://info.arxiv.org/help/api/user-manual.html. Accessed 2026-09-06.

[5] Brave Search. Brave Search API Documentation, 2026. URL https://api-dashboard.search.brave.com/app/documentation/web-search/get-started. Accessed 2026-09-06.

[6] Citation File Format developers. Citation file format 1.2.0, 2021. URL https://citation-file-format.github.io/.

[7] Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daume, and Kate Crawford. Datasheets for datasets. *Communications of the ACM*, 64(12):86–92, 2021. doi: 10.1145/3458723.

[8] GitHub Docs. REST API Endpoints for Releases, 2026. URL https://docs.github.com/en/rest/releases/releases. Accessed 2026-09-06.

[9] GitHub Docs. REST API Endpoints for Search, 2026. URL https://docs.github.com/en/rest/search/search. Accessed 2026-09-06.

[10] Google DeepMind. Gemini 3.1 Pro model card, 2026. URL https://deepmind.google/models/model-cards/gemini-3-1-pro/.

[11] Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. *International Conference on Learning Representations*, 2021. URL https://arxiv.org/abs/2009.03300.

[12] Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. *arXiv preprint arXiv:2103.03874*, 2021. URL https://arxiv.org/abs/2103.03874.

[13] Hugging Face. Hugging Face Hub Documentation, 2026. URL https://huggingface.co/docs/hub/index. Accessed 2026-09-06.

[14] Pranab Islam, Anand Kannappan, Douwe Kiela, Rebecca Qian, Nino Scherrer, and Bertie Vidgen. FinanceBench: A new benchmark for financial question answering. *arXiv preprint arXiv:2311.11944*, 2023. URL https://arxiv.org/abs/2311.11944.

[15] Naman Jain, King Han, Alex Gu, Wen-Ding Li, Fanjia Yan, Tianjun Zhang, Sida Wang, Armando Solar-Lezama, Koushik Sen, and Ion Stoica. LiveCodeBench: Holistic and contamination free evaluation of large language models for code. *arXiv preprint arXiv:2403.07974*, 2024. URL https://arxiv.org/abs/2403.07974.

[16] Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? *arXiv preprint arXiv:2310.06770*, 2023. URL https://arxiv.org/abs/2310.06770.

[17] Kaggle. Kaggle API Documentation, 2026. URL https://www.kaggle.com/docs/api. Accessed 2026-09-06.

[18] Daniel S. Katz et al. Recognizing the value of software: a software citation guide. *F1000Research*, 9:1257, 2021. doi: 10.12688/f1000research.26932.2. URL https://doi.org/10.12688/f1000research.26932.2.

[19] Rodney Michael Kinney, Chloe Anastasiades, Russell Authur, Iz Beltagy, Jonathan Bragg, Alexandra Buraczynski, Isabel Cachola, Stefan Candra, Yoganand Chandrasekhar, Arman Cohan, et al. The semantic scholar open data platform. *arXiv preprint arXiv:2301.10140*, 2023. URL https://arxiv.org/abs/2301.10140.

[20] Jon M. Laurent, Joseph D. Janizek, Michael Ruzo, Michaela M. Hinks, Michael J. Hammerling, Siddharth Narayanan, Manvitha Ponnapati, Andrew D. White, and Samuel G. Rodriques. LAB-Bench: Measuring capabilities of language models for biology research. *arXiv preprint arXiv:2407.10362*, 2024. URL https://arxiv.org/abs/2407.10362.

[21] Zongxia Li, Zhongzhi Li, Yucheng Shi, Ruhan Wang, Junyao Yang, Zhichao Liu, Xiyang Wu, Anhao Li, Yue Yu, Ninghao Liu, et al. Long-horizon-terminal-bench: Testing the limits of agents on long-horizon terminal tasks with dense reward-based grading. *arXiv preprint arXiv:2607.08964*, 2026.

[22] LLM Stats. AI and LLM benchmarks 2026: Rankings, scores and results, 2026. URL https://llm-stats.com/benchmarks/. Accessed 2026-09-06.

[23] Mathematical Association of America. American Invitational Mathematics Examination, 2026. URL https://maa.org/math-competitions/american-invitational-mathematics-examination-aime. Accessed 2026-09-06.

[24] Mike A. Merrill et al. Terminal-Bench: Benchmarking agents on hard, realistic tasks in command line interfaces, 2026. URL https://arxiv.org/abs/2601.11868.

[25] Adrian Mirza et al. Are large language models superhuman chemists? *Nature Chemistry*, 2025. doi: 10.1038/s41557-025-01788-5. URL https://doi.org/10.1038/s41557-025-01788-5.

[26] Margaret Mitchell, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchinson, Elena Spitzer, Inioluwa Deborah Raji, and Timnit Gebru. Model cards for model reporting. In *Proceedings of the Conference on Fairness, Accountability, and Transparency*, pages 220–229, 2019. doi: 10.1145/3287560.3287596.

[27] Moonshot AI. Kimi K3 model card, 2026. URL https://huggingface.co/moonshotai/Kimi-K3.

[28] OpenCompass Contributors. Dataset statistics, 2026. URL https://opencompass.readthedocs.io/en/latest/dataset_statistics.html. Accessed 2026-09-06.

[29] OpenReview. OpenReview Documentation, 2026. URL https://docs.openreview.net/. Accessed 2026-09-06.

[30] Long Phan et al. Humanity's last exam. *arXiv preprint arXiv:2501.14249*, 2025. URL https://arxiv.org/abs/2501.14249.

[31] Jason Priem, Heather Piwowar, and Richard Orr. OpenAlex: A fully-open index of scholarly works, authors, venues, institutions, and concepts. *arXiv preprint arXiv:2205.01833*, 2022. URL https://arxiv.org/abs/2205.01833.

[32] Qwen Team. Qwen3 technical report, 2025. URL https://arxiv.org/abs/2505.09388.

[33] Nils Reimers and Iryna Gurevych. Sentence-BERT: Sentence embeddings using Siamese BERT-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*, pages 3982–3992, 2019. doi: 10.18653/v1/D19-1410.

[34] David Rein, Betty Li Hou, Asa Cooper Stickland, Jackson Petty, Richard Yuanzhe Pang, Julien Dirani, Julian Michael, and Samuel R. Bowman. GPQA: A graduate-level google-proof Q&A benchmark. *arXiv preprint arXiv:2311.12022*, 2023. URL https://arxiv.org/abs/2311.12022.

[35] Arfon M. Smith, Daniel S. Katz, and Kyle E. Niemeyer. Software citation principles. *PeerJ Computer Science*, 2:e86, 2016. doi: 10.7717/peerj-cs.86. URL https://doi.org/10.7717/peerj-cs.86.

[36] Xiaoxuan Wang, Ziniu Hu, Pan Lu, Yanqiao Zhu, Jieyu Zhang, Satyen Subramaniam, Arjun R. Loomba, Shichang Zhang, Yizhou Sun, and Wei Wang. SciBench: Evaluating college-level scientific problem-solving abilities of large language models. *arXiv preprint arXiv:2307.10635*, 2023. URL https://arxiv.org/abs/2307.10635.

[37] Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, et al. MMLU-Pro: A more robust and challenging multi-task language understanding benchmark, 2024. URL https://arxiv.org/abs/2406.01574.

[38] Jason Wei, Zhiqing Sun, Spencer Papay, Scott McKinney, Jeffrey Han, Isa Fulford, Hyung Won Chung, Alex Tachard Passos, William Fedus, and Amelia Glaese. BrowseComp: A simple yet challenging benchmark for browsing agents, 2025. URL https://arxiv.org/abs/2504.12516.

[39] L. Xiaopai. Builderpulse: Ai-powered daily intelligence for indie hackers and builders. GitHub, 2026. URL https://github.com/BuilderPulse/BuilderPulse.

[40] Z.ai. GLM-5.2 model card, 2026. URL https://huggingface.co/zai-org/GLM-5.2.

[41] Z.ai. GLM-5 model card, 2026. URL https://huggingface.co/zai-org/GLM-5.

[42] Zenodo. Zenodo REST API Documentation, 2026. URL https://developers.zenodo.org/. Accessed 2026-09-06.

[43] Junwei Zhou, Zhen Sun, Binyu Li, Jiangyu Zhou, Yuexi Pan, Hengyu Wang, Honghe Ren, Xiaohan Jia, Xueyang Zhou, Xiaoyu Cao, et al. Asi-bench: At the dawn of artificial superintelligence. *arXiv preprint arXiv:2608.17271*, 2026.

Appendix

A Source Inventory and Health

Table 6 lists source connectors that completed successfully at the cutoff. Row counts are measured before duplicate removal and ranking. A connector can be healthy with zero eligible rows if its request succeeds and the response can be parsed. Core sources cover papers, shared model or dataset artifacts, and code; all must be healthy before publication. Failures in optional sources are reported but do not block publication.

Connector	Role	Cutoff status	Core
arXiv	Primary papers via selected AI, language, vision, and software categories	Healthy; 0 rows	Yes
Hugging Face Hub	Datasets and Spaces connected to benchmark artifacts	Healthy; 92 rows	Yes
GitHub Search	Code repositories and benchmark artifacts	Healthy; 300 rows; at cap	Yes
GitHub Organizations	Reviewed organization repositories	Healthy; 4 rows	No
Hugging Face Papers	Community-surfaced papers	Healthy; 0 rows	No
Kaggle Datasets	Public benchmark datasets	Healthy; 32 rows	No
Zenodo	DOI-bearing research artifacts	Healthy; 57 rows	No
OpenReview	Conference submissions	Healthy; no eligible rows	No
GitHub Releases	Curated first-party releases	Healthy; 6 rows	No
OpenAlex	Scholarly discovery metadata	Healthy; 200 rows	No

Table 6: Healthy direct discovery connectors at the current cutoff.

Table 7 lists research and engineering feeds operated by the institutions being monitored. They can surface benchmark papers, datasets, model announcements, and engineering posts. The feed collector yielded one relevant record at the cutoff. Some healthy feeds produced zero matches after relevance filtering. Searches restricted to an organization's website cover institutions that lack a verified feed.

Feeds 1-12	Feeds 13-24
Meituan Engineering	Ai2
OpenAI News	Together AI
Google AI	Sakana AI
Google DeepMind	Qwen
Google Research	Ollama
Apple Machine Learning Research	Stability AI
AWS Machine Learning	Nomic AI
Hugging Face Blog	Replicate
Microsoft Research	NVIDIA Developer
NVIDIA AI Blog	IBM Research
Mistral AI	Databricks
Meta Research	LangChain

Table 7: First-party research and engineering feeds monitored by Benchmark Radar.

B Worked Use Case: Prior-Art Check for a New Evaluation

The contributor wanted to determine whether a proposed evaluation duplicated existing work. During August 2026, they surveyed recent work on credit assignment in agentic training, using small Qwen-series models as a requirement for reproducible baseline experiments. The search produced a focused comparison table that the contributor used to assess the proposed design.

The contributor gave the task to a coding agent together with the public setup instructions for Benchmark Radar. The agent installed the command-line client and the Benchmark Radar Skill, a set of instructions for using the client. It then downloaded the local corpus and queried candidate records offline.

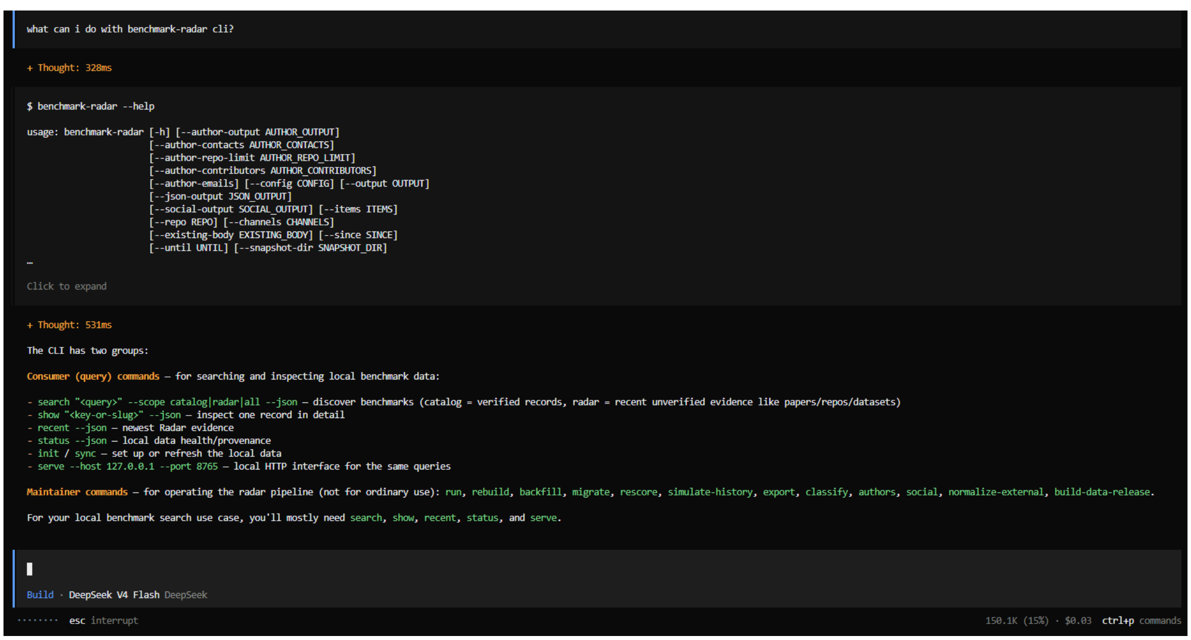

Figure 5: A coding agent installs the Benchmark Radar client and its usage instructions, then runs local queries.

The agent inspected paper, repository, and dataset links to decide which sources to open next. Figures 6 and 7 show two retrieved records: one has a repository link, and the other has both paper and repository links. Neither records a dataset link, leaving dataset availability for a follow-up check.

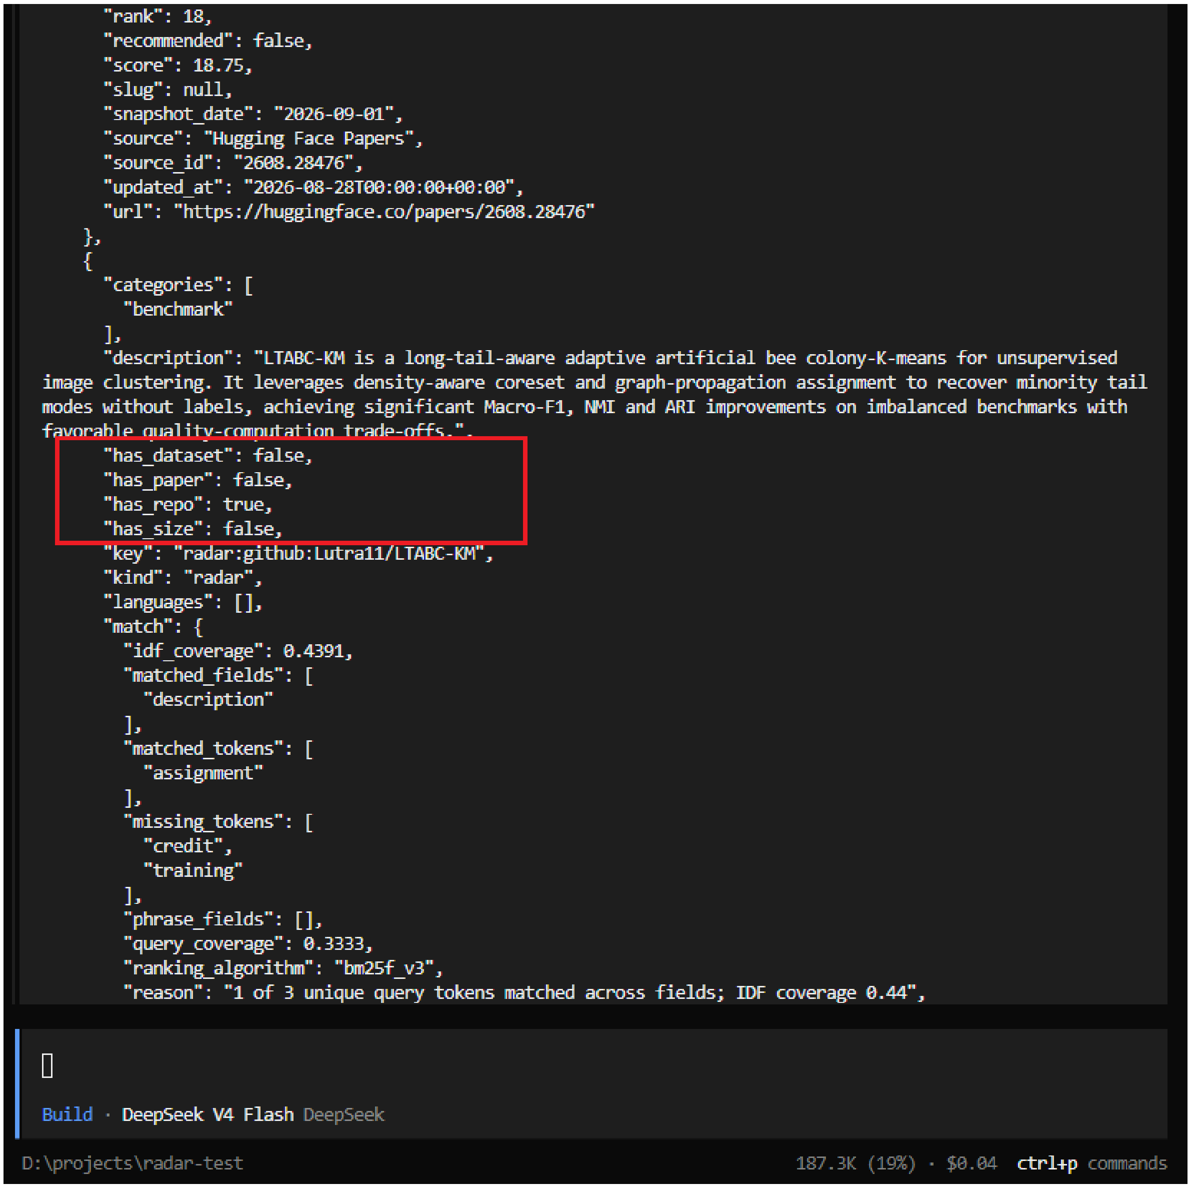

```
      "rank": 18,
      "recommended": false,
      "score": 18.75,
      "slug": null,
      "snapshot_date": "2026-09-01",
      "source": "Hugging Face Papers",
      "source_id": "2608.28476",
      "updated_at": "2026-08-28T00:00:00+00:00",
      "url": "https://huggingface.co/papers/2608.28476"
    },
    {
      "categories": [
        "benchmark"
      ],
      "description": "LTABC-KM is a long-tail-aware adaptive artificial bee colony-K-means for unsupervised
image clustering. It leverages density-aware coreset and graph-propagation assignment to recover minority tail
modes without labels, achieving significant Macro-F1, NMI and ARI improvements on imbalanced benchmarks with
favorable quality-computation trade-offs.",
      "has_dataset": false,
      "has_paper": false,
      "has_repo": true,
      "has_size": false,
      "key": "radar:github:Lutra11/LTABC-KM",
      "kind": "radar",
      "languages": [],
      "match": {
        "idf_coverage": 0.4391,
        "matched_fields": [
          "description"
        ],
        "matched_tokens": [
          "assignment"
        ],
        "missing_tokens": [
          "credit",
          "training"
        ],
        "phrase_fields": [],
        "query_coverage": 0.3333,
        "ranking_algorithm": "bm25f_v3",
        "reason": "1 of 3 unique query tokens matched across fields; IDF coverage 0.44",
```


Figure 6: A candidate record with a repository link for inspecting the implementation. Paper and dataset links are missing from the record.

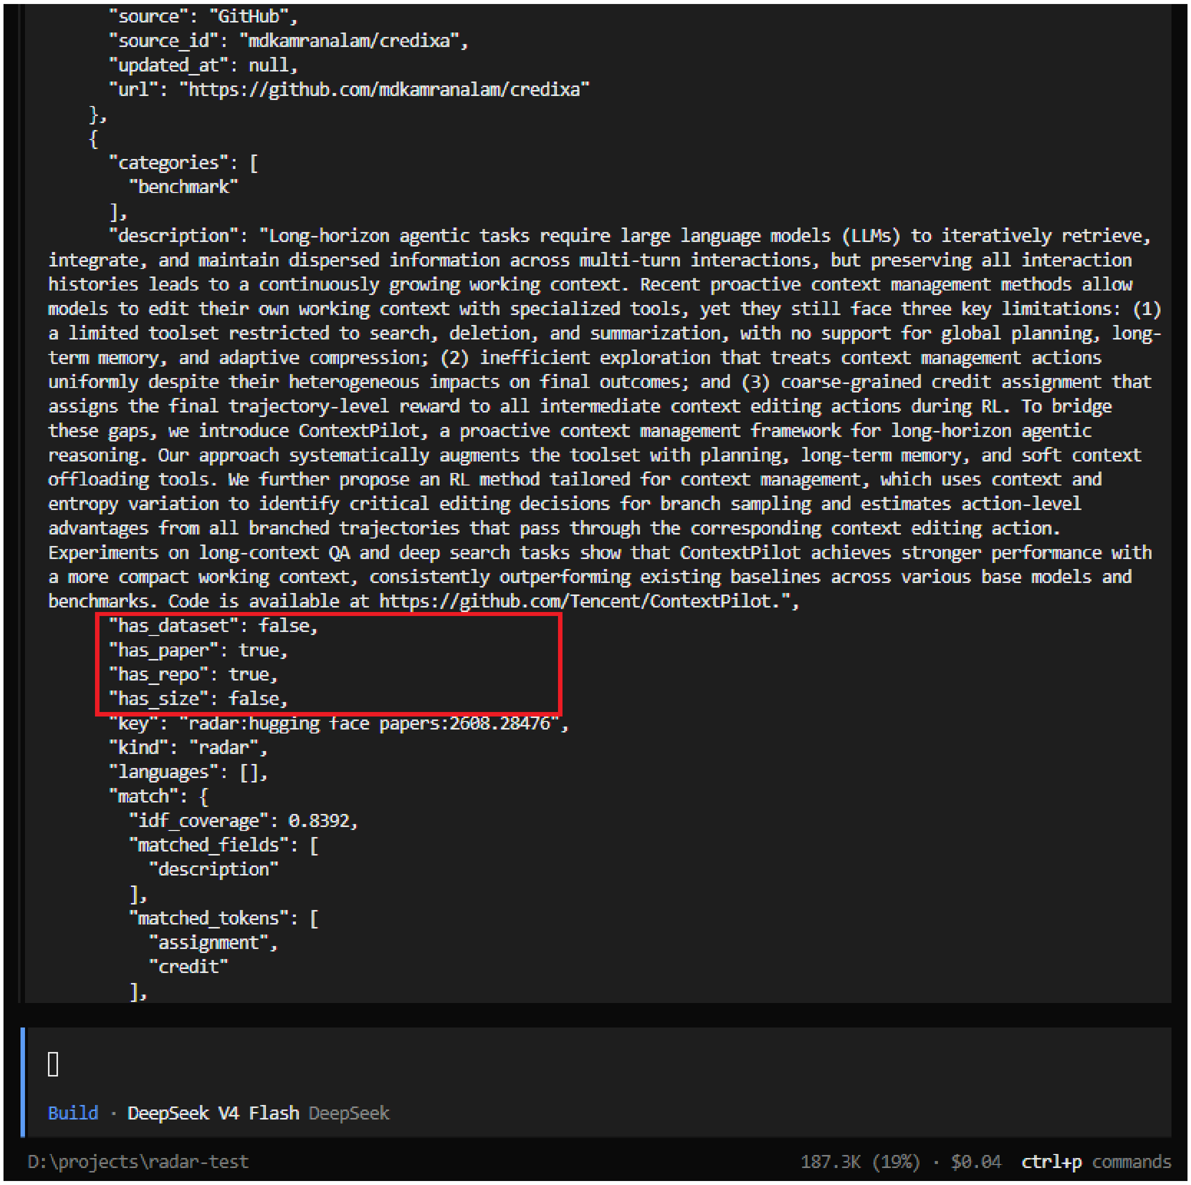

Figure 7: A second candidate provides both paper and repository links. Its dataset link needs a follow-up check.

The agent supplemented catalog queries with web search to find related work described in different terms. It cross-checked the candidate sets and inspected the source evidence before selecting records for the comparison. Figure 8 shows the word matches and score components available for inspecting one retrieved candidate.

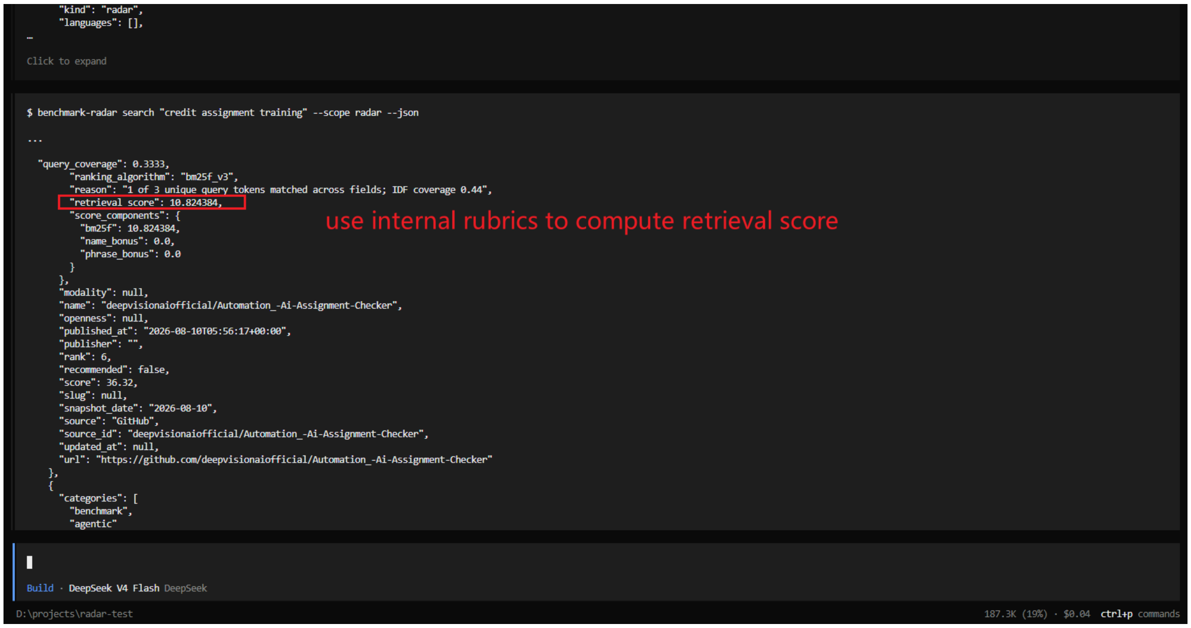

Figure 8: A candidate record with query-word matches and retrieval-score components for the agent to inspect when judging relevance.

Figure 9 shows the resulting related-work table, which the contributor judged complete enough to act on.

#	Work	Published	Qwen base model(s)	Credit-assignment focus	Agentic benchmarks used	Link
1	SRPO: Self-Reflective Policy Optimization for Long-Horizon Reasoning	2026-08-25	Qwen3-8B (also Qwen3-1.7B / 32B scaling)	Reflection-conditioned dense token-level signals convert sparse terminal reward into training signal	AIME'24, WebShop, ALFWorld, SWE-Bench-Lite	https://arxiv.org/abs/2608.23493 · https://github.com/Galleons2029/SRPO
2	ContextPilot: Teaching Agents for Proactive Context Management via Fine-grained RL	2026-08-28	Qwen3-8B, Qwen3-14B	Coarse-grained → action-level credit assignment over branched trajectories during agentic RL	InfBench (InfiniteBench), NovelQA, LongMemEval, BrowseComp+	https://arxiv.org/abs/2608.28476 · https://github.com/Tencent/ContextPilot
3	SkillGate: Training In-Policy Skill Selection in Long-Horizon Agents	2026-08-21	Qwen3.5-9B (SFT checkpoint, RL init)	"Selector credit starvation"; partitions token support into two credit channels (outcome vs. action-local advantage)	5 agentic benchmarks incl. SkillsBench, SWE (16-candidate skill slate)	https://arxiv.org/abs/2608.18852
4	CIPO: Contextual Information Policy Optimization for Search Agents	2026-08-06	Qwen2.5-3B-Instruct, Qwen2.5-7B-Instruct	Dense turn-level credit (EALR) to evidence-using reasoning actions, combined with outcome reward	HotpotQA, 2WikiMultiHopQA, MuSiQue, Bamboogle + 3 OOD	https://arxiv.org/abs/2608.06128
5	MoRSE: Task-Oriented Multi-Agent System with Mixture of Role-Subtask Experts	2026-08-10	Qwen3-4B-Instruct (one of three backbones; also Llama-3.1-8B, Gemma-4-31B)	Hierarchical GRPO with two-layer credit assignment (isolates expert vs. routing quality)	Code-generation benchmarks (multi-agent), held-out task domains	https://arxiv.org/abs/2608.09251

Figure 9: Summary table of recent work on credit assignment in agentic training, assembled during the session.

Figure 10 shows a comparison table from an earlier benchmark-design effort by one contributor. It covers ten related benchmarks across six design dimensions, with each cell read from a source paper. The contributor reported that assembling it took more effort than any part of that project except producing the benchmark data. Benchmark Radar retrieves candidates for this kind of comparison and exposes the matching words and fields for inspection; the comparison itself requires reading the source evidence.

Bench Name	End-to-End Tasks	Fine-Grained Eval	Researcher-Quality Eval	Data Generation	Multi-Harness Eval	#Tasks
MLE-Bench (Chan et al., 2025)	✗	✓	✗	Transfer&Compose	✓	75
MLGym-Bench (Nathani et al., 2025)	✗	✓	✗	Automatic	✗	13
EXP-Bench (Kon et al., 2025)	✓	✗	✗	Automatic	✓	461
ResearchCodeBench (Hua et al., 2026)	✗	✓	✗	Transfer&Compose	✗	212
MLR-Bench (Chen et al., 2026)	✓	✗	✗	Automatic	✓	201
PaperBench (Starace et al., 2025)	✗	✓	✗	Transfer&Compose	✗	8316
AstaBench (Bragg et al., 2025)	✓	✓	✗	Transfer&Compose	✓	2400+
InnovatorBench (Wu et al., 2025)	✓	✗	✗	Transfer&Compose	✗	20
AIRS-Bench (Lupidi et al., 2026)	✗	✓	✗	Automatic	✓	20
COMPOSITE-Stem (Waters et al., 2026)	✓	✓	✗	Manual	✗	70
ScienceBoard (Sun et al., 2025)	✗	✓	✗	Manual	✗	169

Figure 10: A prior-art comparison table assembled manually for an earlier benchmark-design effort. Each row is one related benchmark and each column one design dimension, all read from the source papers by hand.

This case records one workflow combining Benchmark Radar with web search; it has no measured time or accuracy comparison.[3] The summary table and session evidence are available in the project's public issue #492.

C Reproducibility, Access, and Citation

This paper evaluates Benchmark Radar v0.10.0 at Git commit 76a6180 with a data cutoff of 2026-09-06. The data build combines dated collection snapshots, benchmark registry records, and model reports, then classifies tracks and packages the results with checksums. Installed clients validate the checksum and record format before activating an update. The evaluation used a clean data rebuild with linting, formatting checks, catalog normalization, classification, release construction, and the full test suite; all 1,236 tests passed. The manuscript is authored in LaTeX, separately from the data generators. The project provides a DOI and a Citation File Format record for citation [6, 18, 35].

Artifact	Canonical or permanent location
Archived paper (v0.9.0)	`https://doi.org/10.5281/zenodo.22167102`
Source code	`https://github.com/ktwu01/benchmark-radar`
Dashboard	`https://benchmark-radar.org/`
Discovery observations (JSON)	`https://benchmark-radar.org/data/radar.json`
Benchmark catalog	`https://benchmark-radar.org/data/benchmark-index.json`
RSS	`https://benchmark-radar.org/feed.xml`
Citation metadata	`https://github.com/ktwu01/benchmark-radar/blob/main/CITATION.cff`

Table 8: Access and citation artifacts for Benchmark Radar.

The reported counts were computed from the cumulative discovery data, benchmark catalog, model registry, model-report and score archives, normalized catalog records, and collection configuration. Their repository locations are `site/data/radar.json`, `site/data/benchmark-index.json`, `site/data/models.json`, `data/model_cards.yml`, `data/benchmark_scores.yml`, `data/catalog/`, and `config.yml`. The dated numerical inputs used by the manuscript and figures are stored in `figure-data.tex` beside the LaTeX source.

[3] Jiayu Wang contributed this worked use case and its supporting evidence.

Tables 4 and 5 use the model-report and score archives at commit 76a6180. For the coverage gaps, take each benchmark's latest score date and latest model-report mention date, then retain gaps of at least 180 days. For headroom, select the highest numeric score in each model-report track with a percentage scale and retain values at least 95. Each selected observation retains its model, source, instrument, and protocol.

To reproduce the data build from a clean checkout, install the project and run `benchmark-radar normalize-catalog`, `benchmark-radar classify`, and `benchmark-radar build-data-release`, in that order. Run `make` in `docs/technical-report/latex/` to build this manuscript from its dated numerical inputs. Rebuilding the data at a later revision does not update those inputs automatically.

The DOI in Table 8 identifies the archived v0.9.0 paper. This manuscript describes v0.10.0 at the cutoff above. The dashboard continues to change; citations of its live counts need a retrieval date.